

MUNEEB RASHID

Senior AI Engineer | Microsoft Certified
Generative AI · Agentic AI · MLOps & LLMOps
Riyadh, Saudi Arabia · muneebcyan@gmail.com
+92 333 076 0460 · Iqama Status: Transferable
[LinkedIn](#) | [GitHub](#) | [Portfolio](#)

PROFESSIONAL SUMMARY

Senior AI Engineer with **5+ years** turning complex business requirements into production AI systems across healthcare, multimedia, and workflow automation. Designed RAG and agent workflows for platforms serving **100,000+ users**, with ownership spanning architecture, implementation, evaluation, and deployment. Combines **LLM fine-tuning and inference optimization** with Azure MLOps and on-device computer vision. Leads client-facing technical delivery, builds quality gates into model releases, and mentors engineers to establish consistent development practices.

PROFESSIONAL EXPERIENCE

- Senior AI Engineer

Arbisoft

Riyadh, Saudi Arabia

Nov 2024 – Present

 - Led technical delivery of **Storybook Studio**; designed LangGraph/LlamaIndex orchestration integrating text, image, video, and audio APIs for a platform serving **10,000+ daily active users**
 - Implemented **character-consistency workflows across 50+ scenes** and a Gemini Flash validation layer to identify generation defects; integrated voiceover across **15+ languages**
 - Architected **AgentOps**: MLflow tracking and registry, LangGraph services on Azure Container Apps, and Azure DevOps CI/CD; integrated **quality gates and rollback** into the model-release workflow
 - Implemented **FlowAssist** conversational agents and n8n/Zapier workflow integrations; packaged services with Docker and Azure Container Registry for deployment
- Senior AI Engineer

PureLogics

Remote, USA

May 2023 – Oct 2024

 - Designed and implemented the **MedQuery RAG pipeline** within a healthcare platform serving **100,000+ active users**; added PII redaction, access controls, and audit logging to the AI workflow
 - Engineered **DataLens** multimodal RAG on AWS Bedrock; implemented elastic scaling and vector-index optimization to manage concurrent demand and infrastructure cost
 - Worked directly with US and Canadian enterprise clients to scope healthcare, fintech, and e-commerce automation requirements; translated them into AI workflows and delivered production handover
- Machine Learning Engineer

SmartFun Studios

Remote, USA

Jun 2020 – Apr 2023

 - Cut **FaceLite** inference from **8s to 800ms (90% faster)** on Android Mali-T860 mobile GPU through model quantization and hardware-specific optimization, enabling local inference
 - Implemented **DriveGuard** driver-drowsiness detection and real-time alerts, connecting model inference to the operational notification workflow
 - Established the team's **DVC/Git experimentation and model-evaluation workflow**; mentored **three junior engineers** and authored ML training documentation to standardize development practices
- Machine Learning Engineer (Part-time)

Yottabyte

Pakistan

Jun 2018 – Feb 2020

 - Designed Azure ML + Azure DevOps **MLOps pipelines** automating training, evaluation, and deployment
 - Implemented Docker/ACR packaging, automated quality gates, and rollback for classification and regression deployments

TECHNICAL EXPERTISE

- Agents & Retrieval

LangGraph, LangChain, LlamaIndex, MCP, hybrid search, GraphRAG/Neo4j, reranking, Pinecone, Weaviate
- Models & Evaluation

PyTorch, LoRA/QLoRA, DeepSpeed, DPO, vLLM, AWQ, ONNX, TensorRT, OpenCV, TensorFlow, RAGAS, DeepEval, LangSmith, Langfuse
- Cloud & MLOps

Azure ML, AI Search, Container Apps, Azure DevOps, MLflow, DVC, AWS Bedrock, GCP Vertex AI, Docker, Kubernetes
- Backend & Controls

Python, SQL, FastAPI, PostgreSQL, Redis, PII redaction, access controls, audit logging, regression gates, canary deployment

SELECTED ENGINEERING PROJECTS

Distributed LLM Fine-Tuning & Inference Platform

DeepSpeed ZeRO-3 · QLoRA · vLLM · AWQ · MLflow

Built a platform to adapt LLaMA 3 and Mistral models while managing GPU memory and serving latency. Implemented **4-bit QLoRA with DeepSpeed ZeRO-3 sharding**, gradient checkpointing, and optimizer offload to manage training memory across GPUs. Built the adapter merge and AWQ quantization path into **vLLM serving with continuous batching and speculative decoding**. Tracked runs in MLflow and versioned adapters for evaluation-gated release, connecting training artifacts to deployable model versions.

Self-Improving Agent Platform

LangGraph · DPO · LangSmith · Azure ML · GitHub Actions

Designed and implemented a **trace-to-training workflow**: converted agent failures and user corrections into preference pairs for DPO updates to routing and response models. Built an evaluation harness combining deterministic checks, LLM-based scoring, and **bootstrap confidence intervals** to compare candidates before promotion. Integrated canary deployment and rollback on groundedness or task-success regression, making model updates conditional on evaluation outcomes.

Edge Vision Fleet Platform

ONNX · TensorRT INT8 · Jetson · Kubernetes · Prometheus

Built model delivery workflows for an edge vision fleet. Exported models to ONNX and used **TensorRT INT8 on NVIDIA Jetson**, benchmarking accuracy and latency per target before release. Shipped signed, versioned OTA bundles through **1% then 10% canaries before fleet rollout**, with rollback on confidence drift. Kept inference local and escalated low-confidence frames to a cloud model, limiting cloud calls to uncertain cases.

CERTIFICATION & PROFESSIONAL DEVELOPMENT

Microsoft Certified: Machine Learning Operations Engineer Associate (AI-300)

Training: Azure Machine Learning & MLOps; MLOps Bootcamp — Udemy

EDUCATION

MS in Data Science — FAST NUCES, Islamabad

2021–2023

GPA: 3.73/4.0 | Teaching Assistant: Machine Learning & Advanced ML

BS in Mathematics — Bahauddin Zakariya University, Multan

2017–2021

GPA: 3.53/4.0